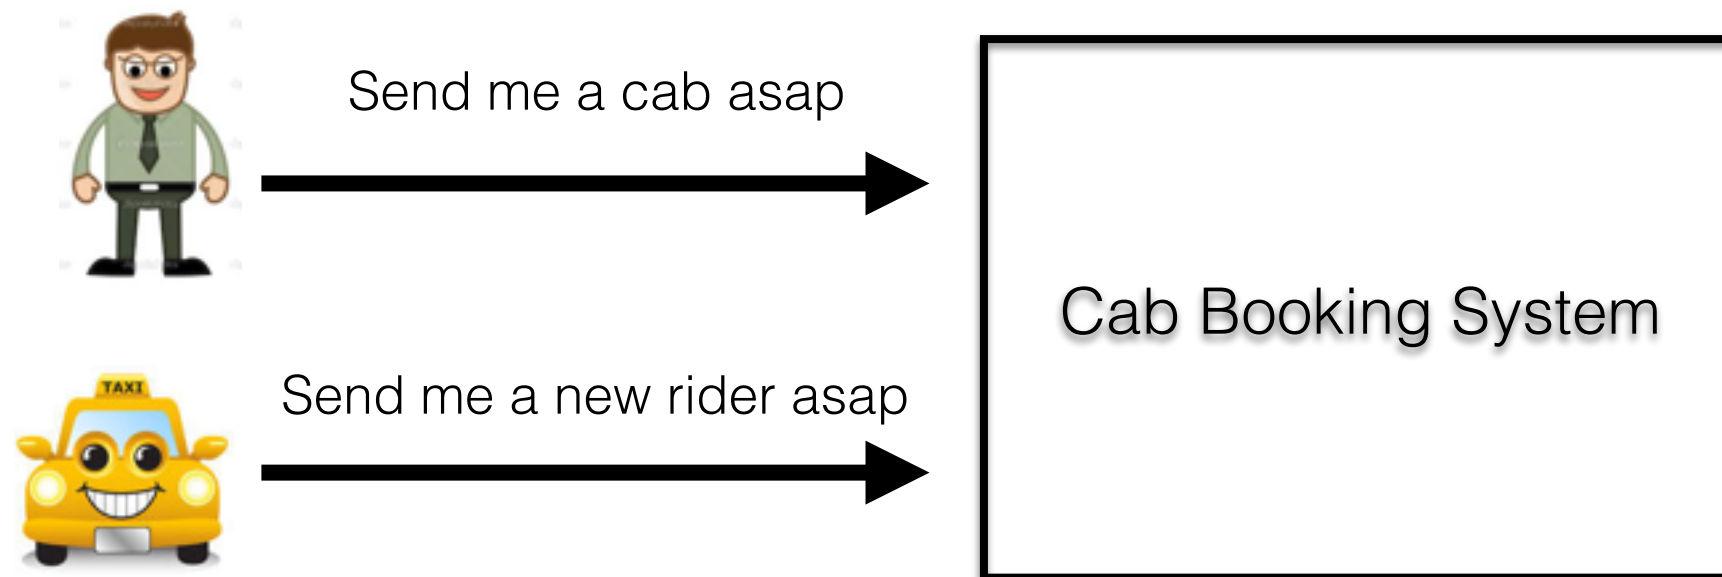


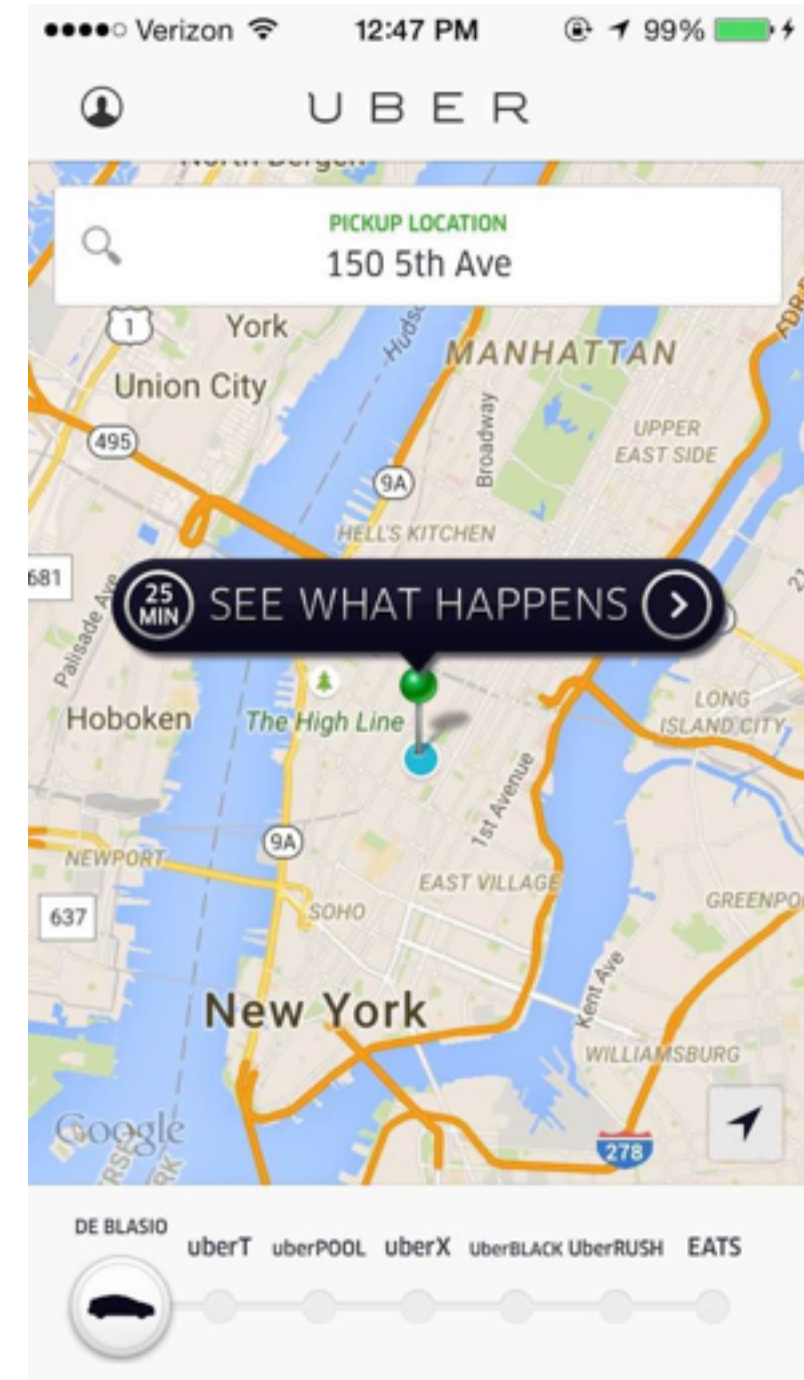
Book My Cab: Hop In, Let's Go!

Umesh Shukla

A Platform to book Cabs



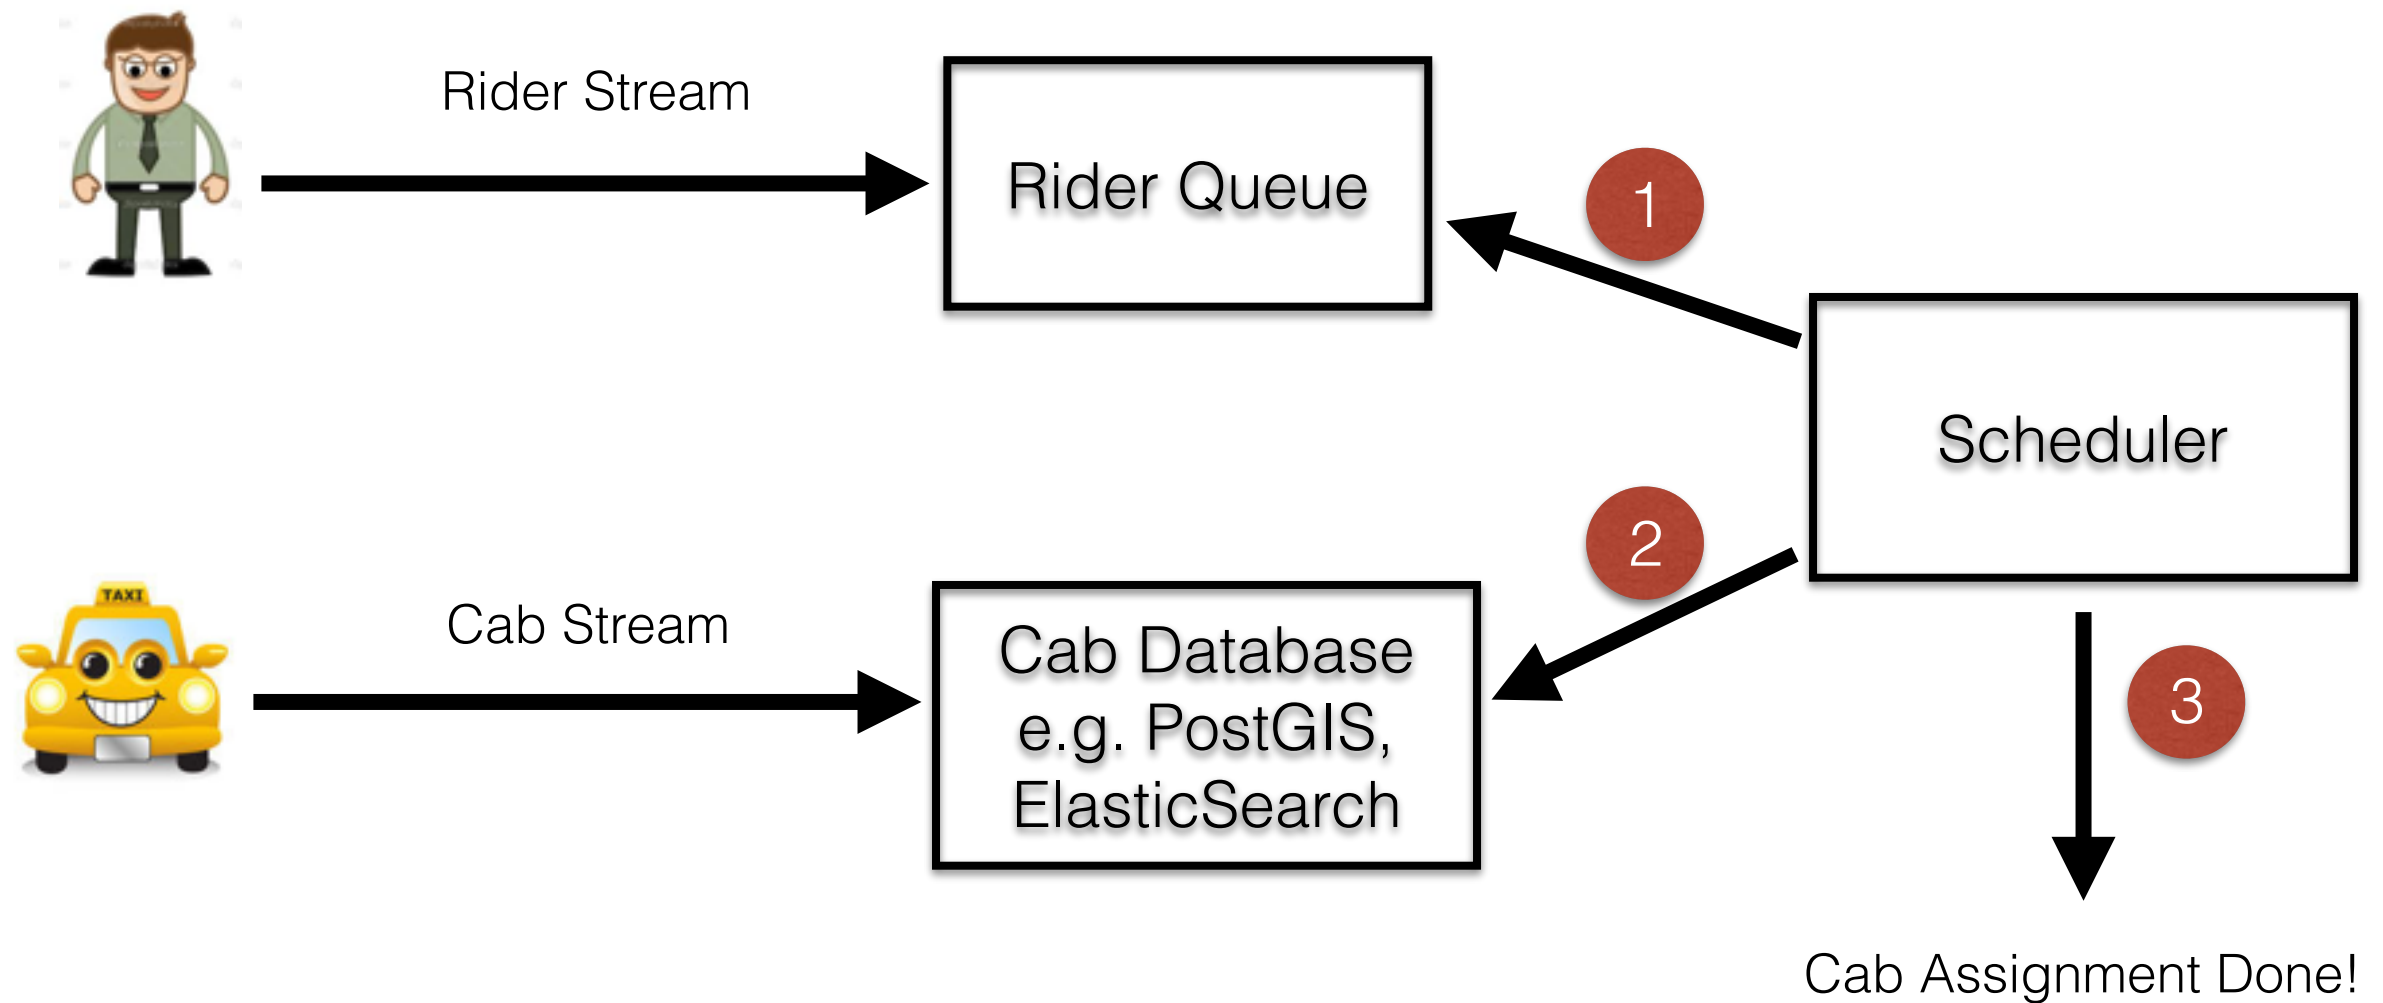
How to do design a scalable real-time cab booking system?



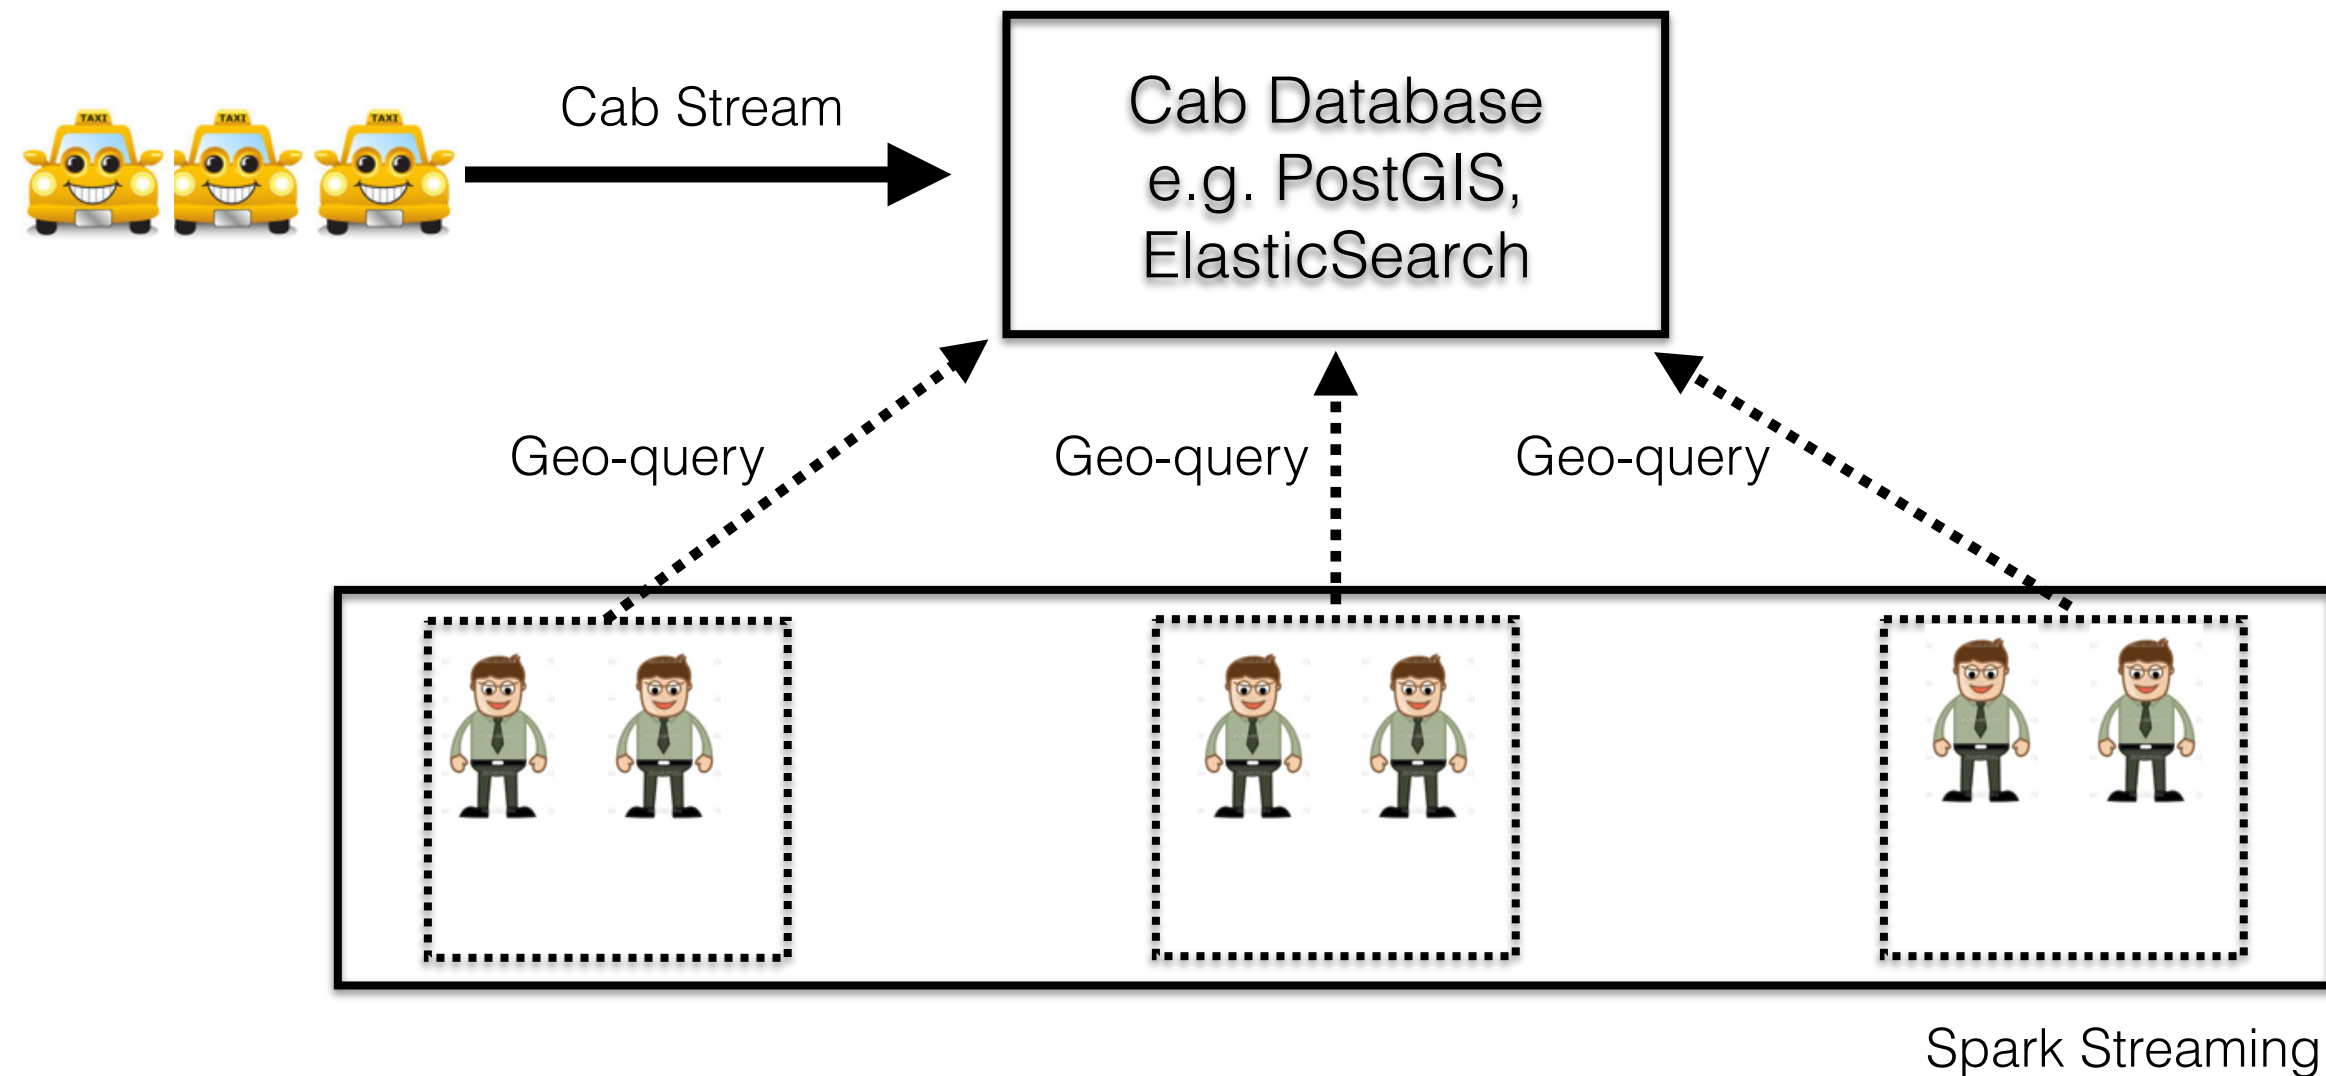
Demo

<http://bit.ly/2dquv0B>
(Might be down)

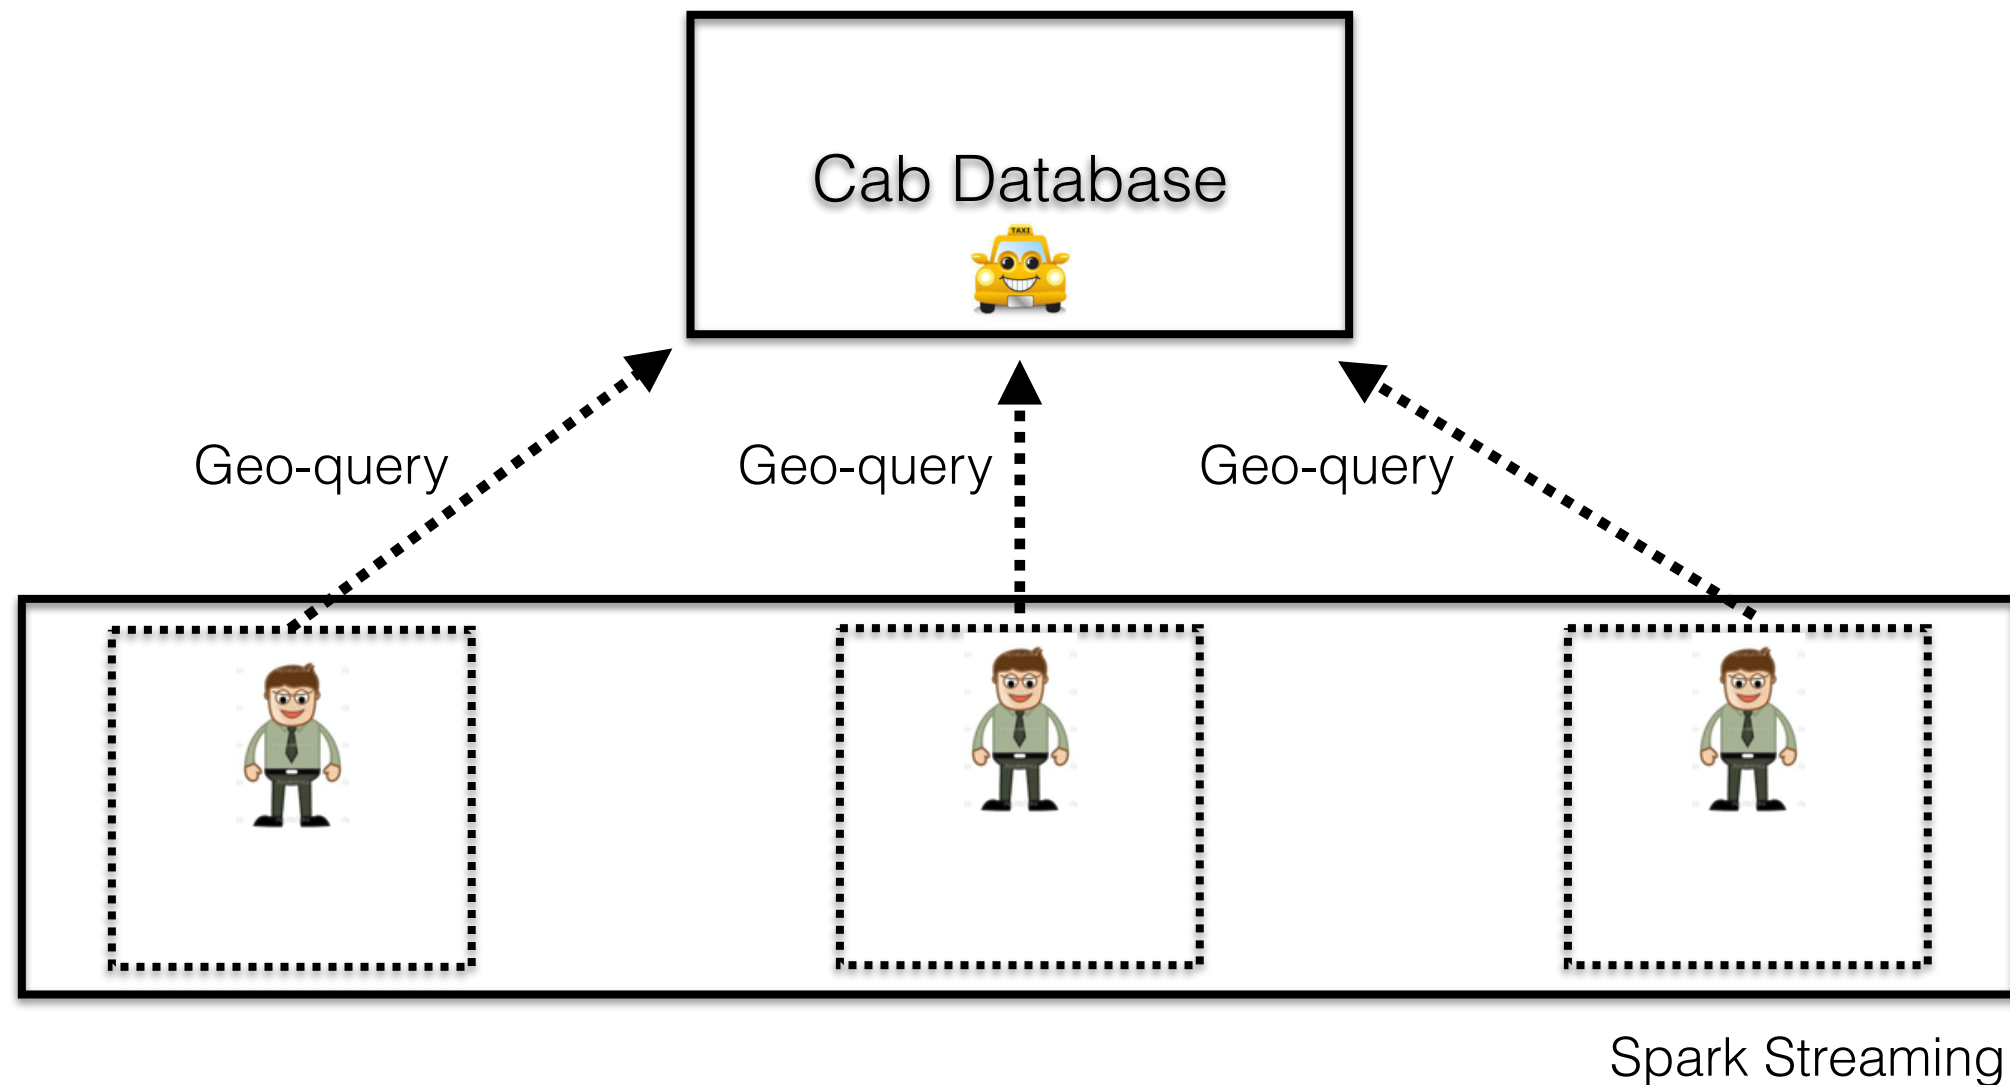
Basic Algorithm



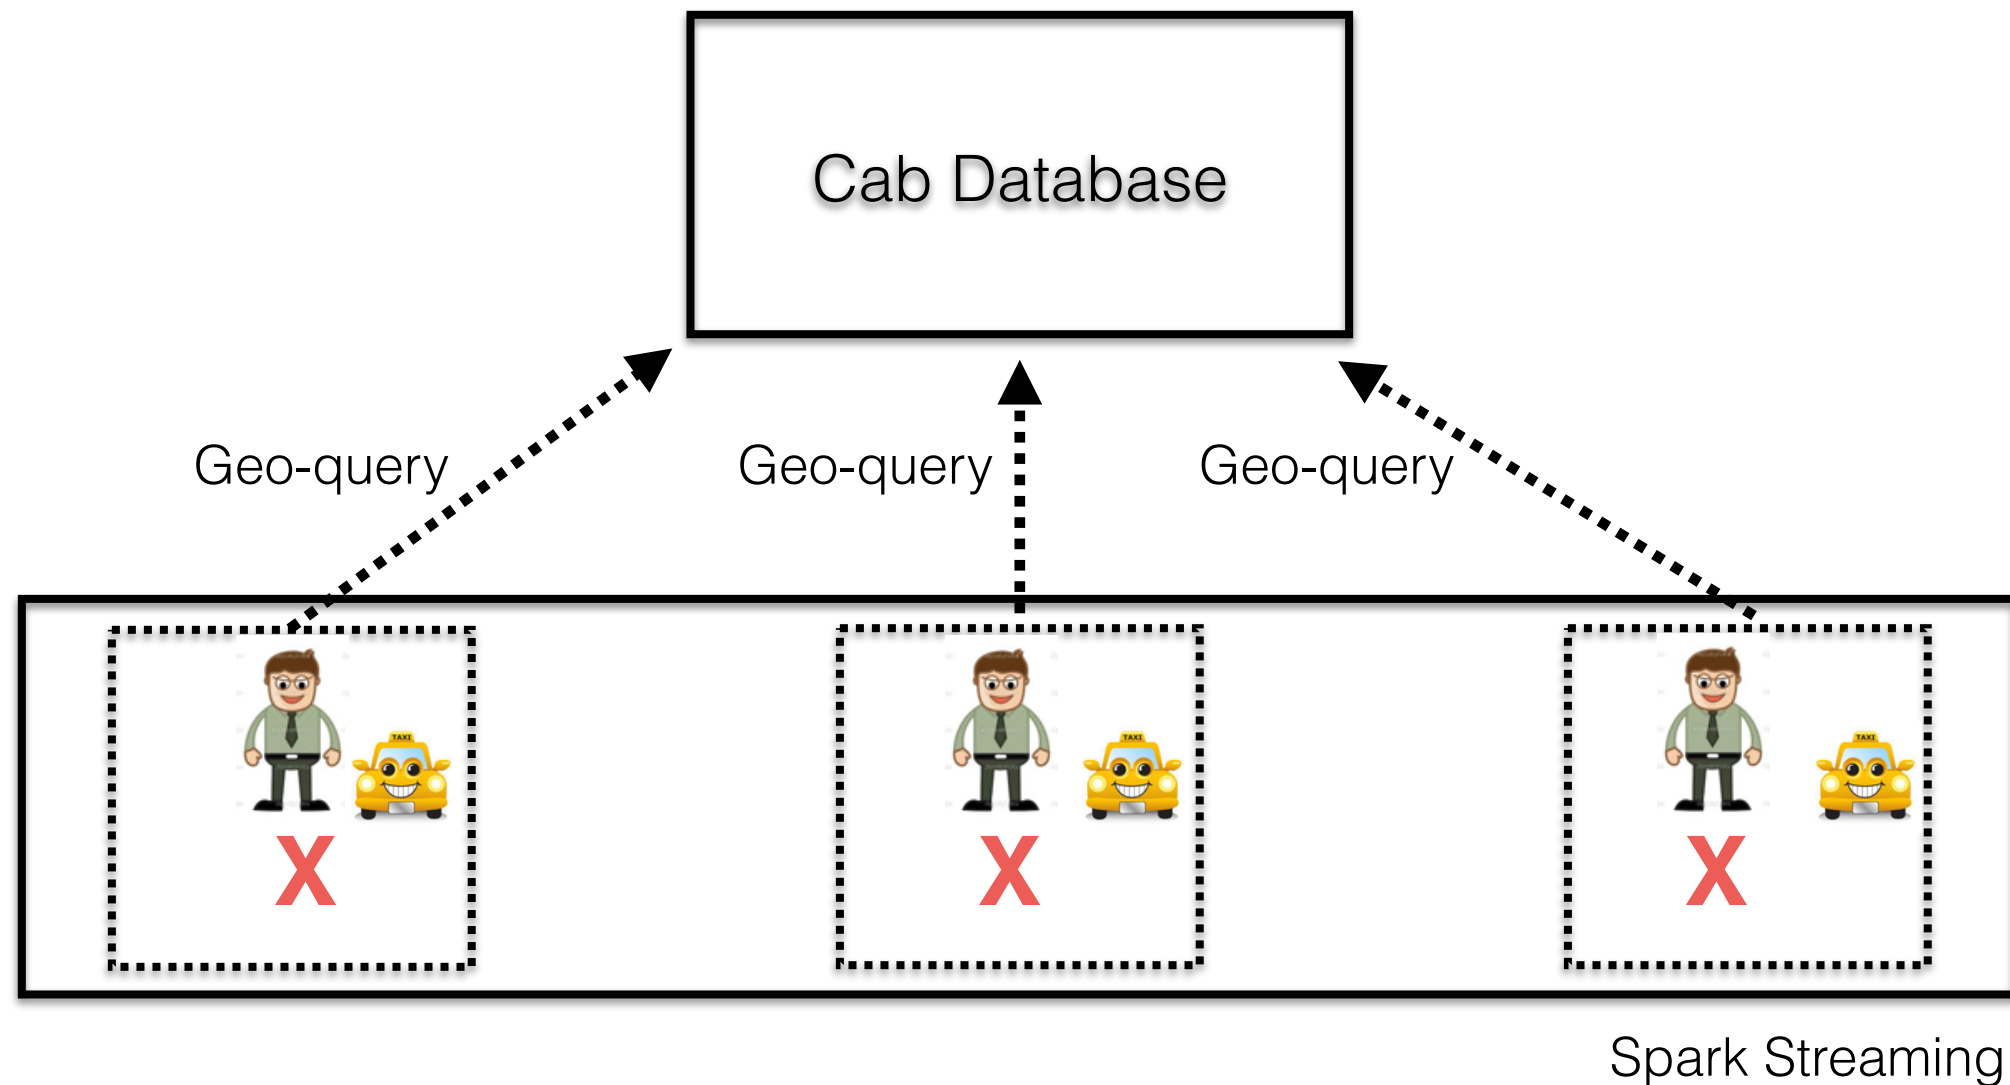
How to Design this System in Distributed & Parallel Fashion?



Over Booking Scenario

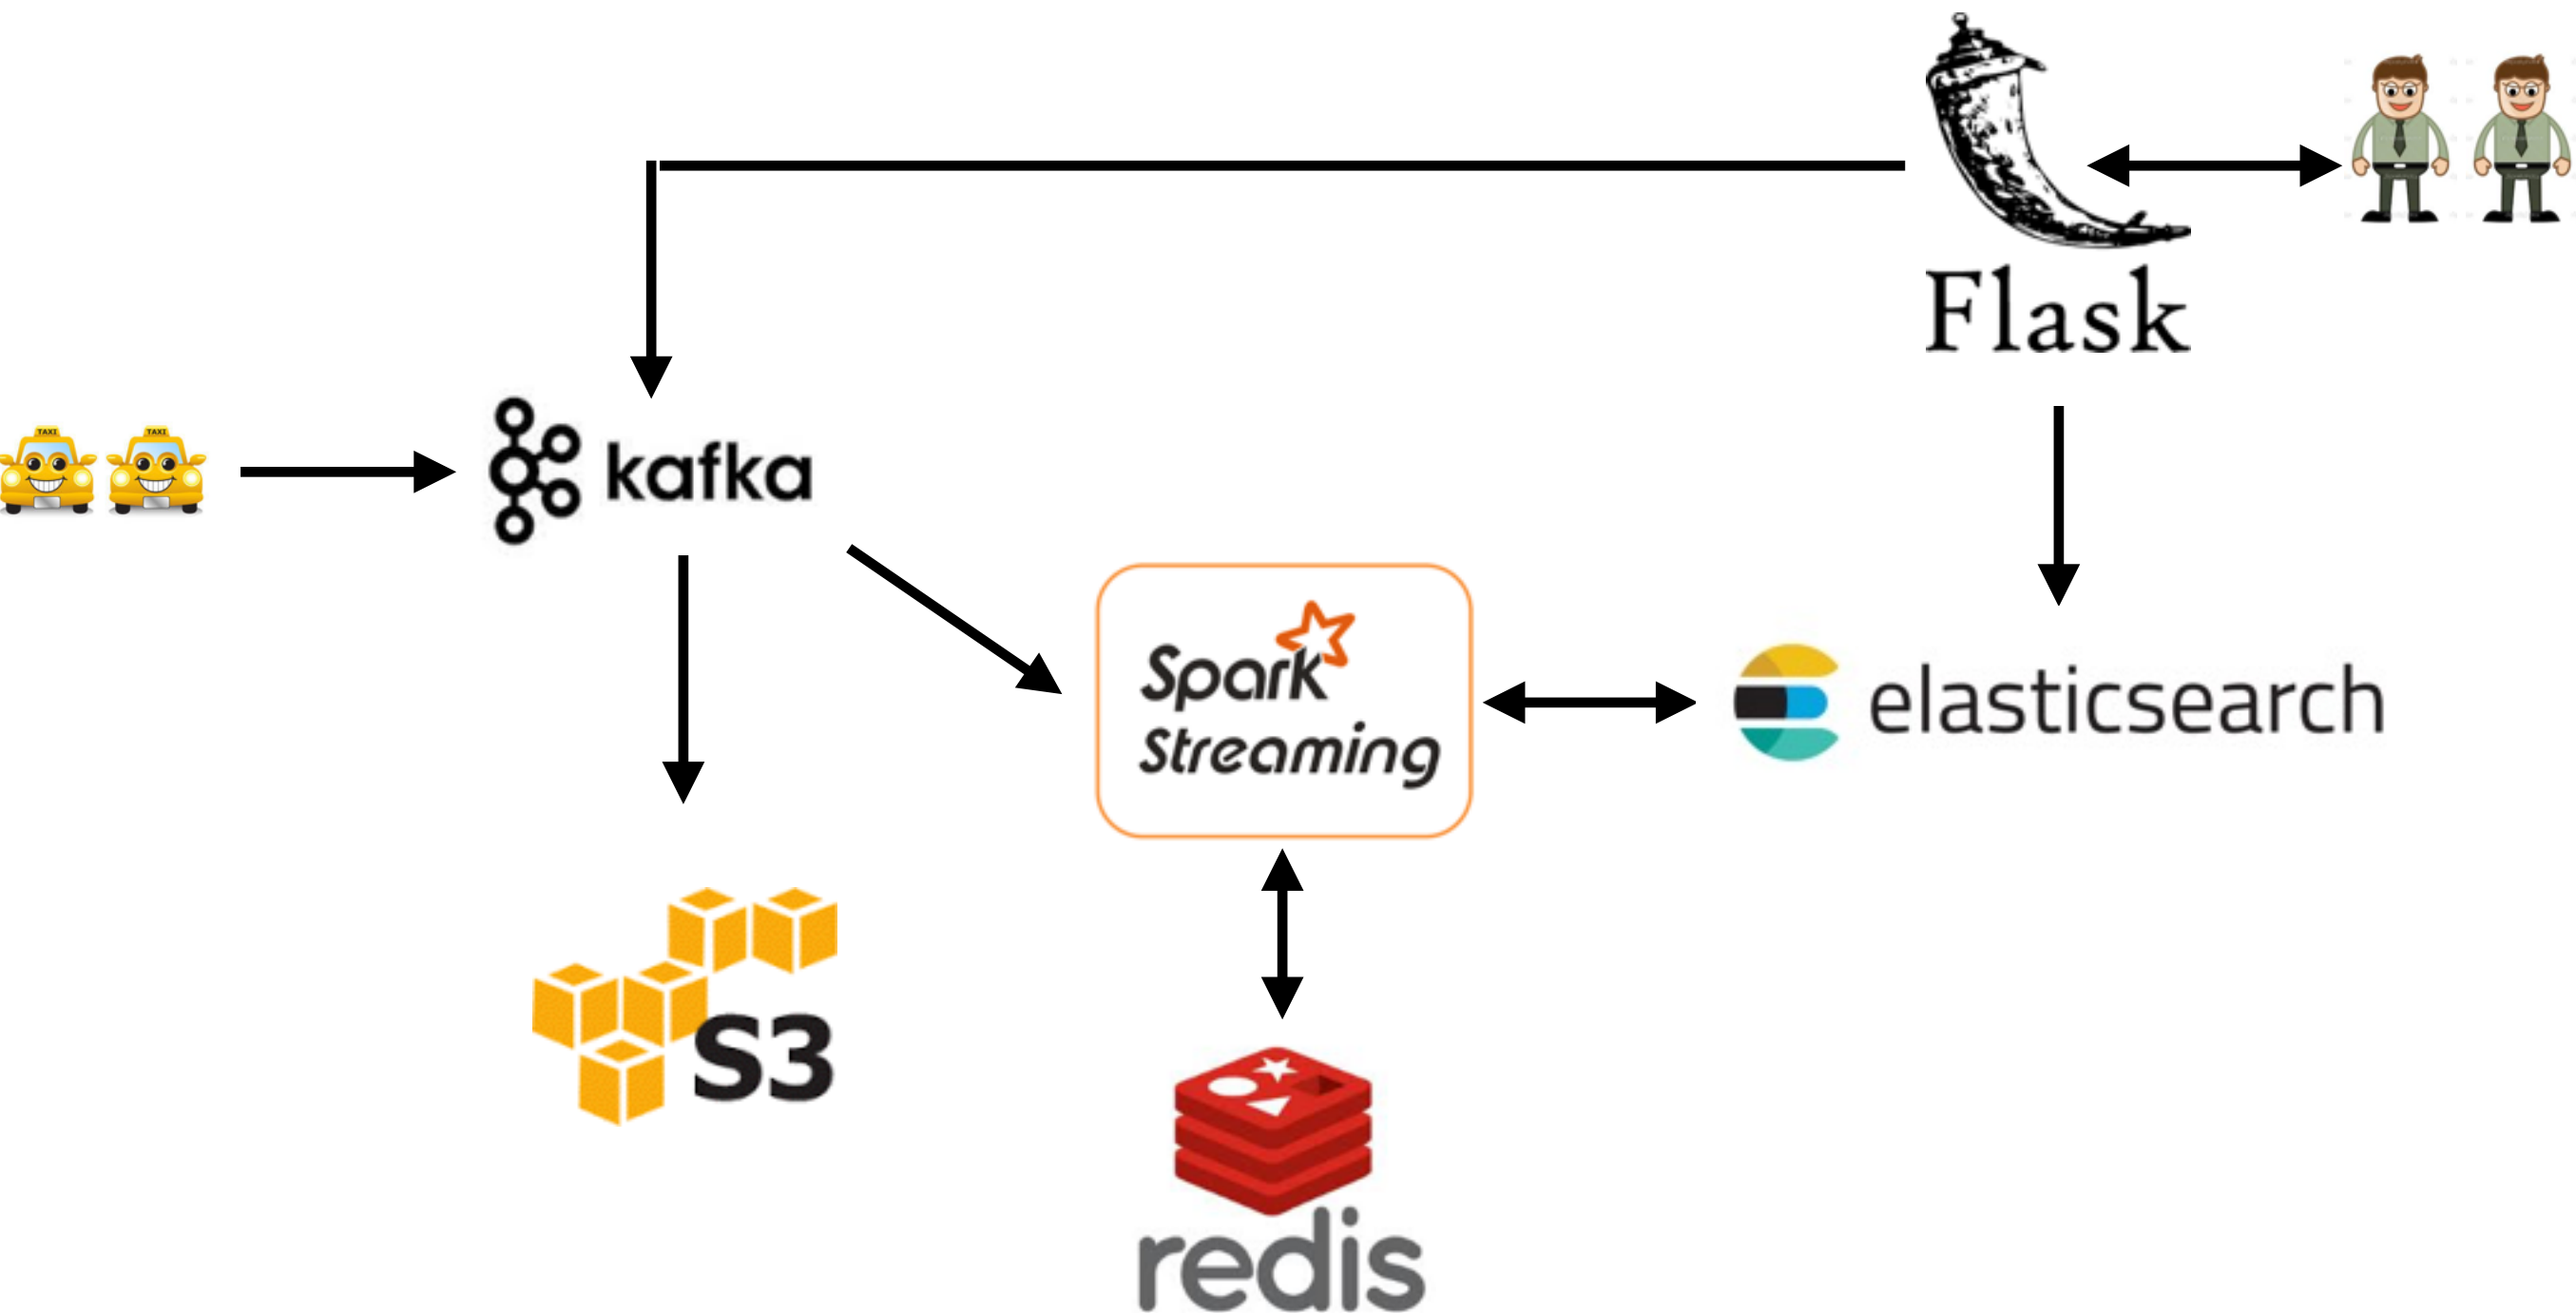


Over Booking Scenario



Solution: Use a transactional Cache e.g. Redis

Data Pipeline



Performance Numbers!

- Playback with NYC Yellow Cab data, 4-node AWS cluster
- System completes assignment of ~ 2000 customers and ~2000 drivers in 4-5 s.
Spark Streaming Window Size: 5 s

Next Steps

- Take the full advantage of a micro-batch to come up with even more optimal cab assignment algorithm to improve cab utilization e.g. Maximal Weight Matching
- Challenge: More computationally intensive $O(N^3)$ vs $O(N^2)$